

# **Detecting Cancer Metastases on Gigapixel Pathology Images**

**COMS W4995 010 Applied Deep Learning**

**Instructor: Joshua Gordon**

**Team Members: Weiwei Zhan, Haokai Zhao**

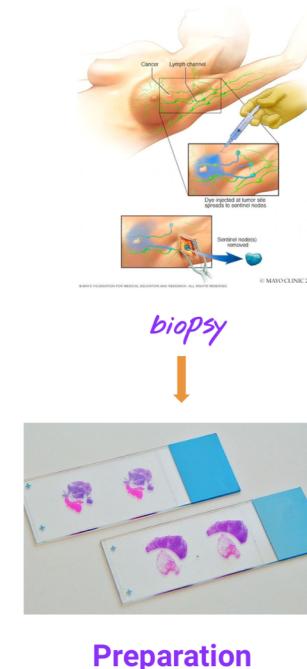
# Background

Microscopic examination of breast cancer

- Time-consuming and error-prone
- Computer assisted detection

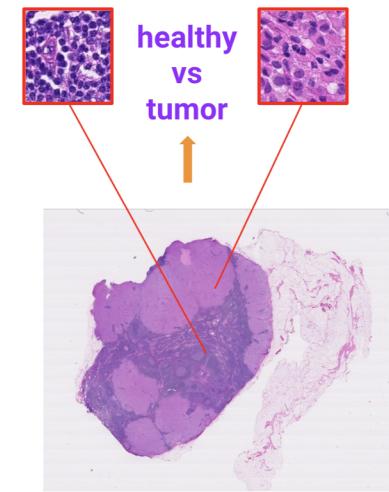
Deep CNNs have significantly improved the accuracy of computer vision tasks

**Goal:** to develop a CNN framework to assist breast cancer detection



Preparation

Diagnosis → Treatment plan



Visual inspection

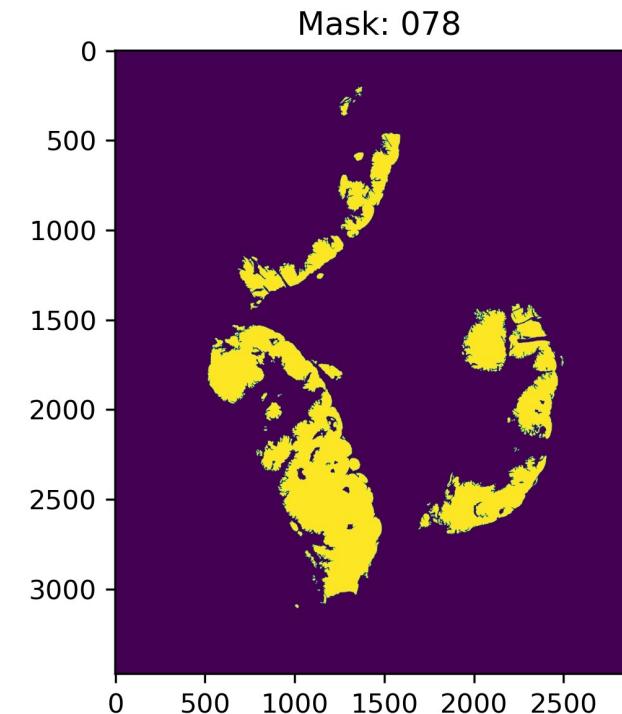
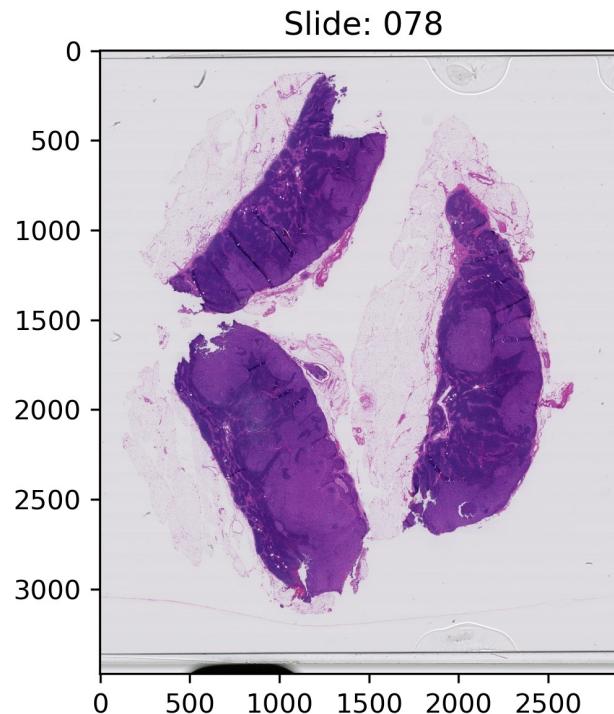
❖ From Course Slides

# I. Preprocessing

# 1.1 Load the dataset

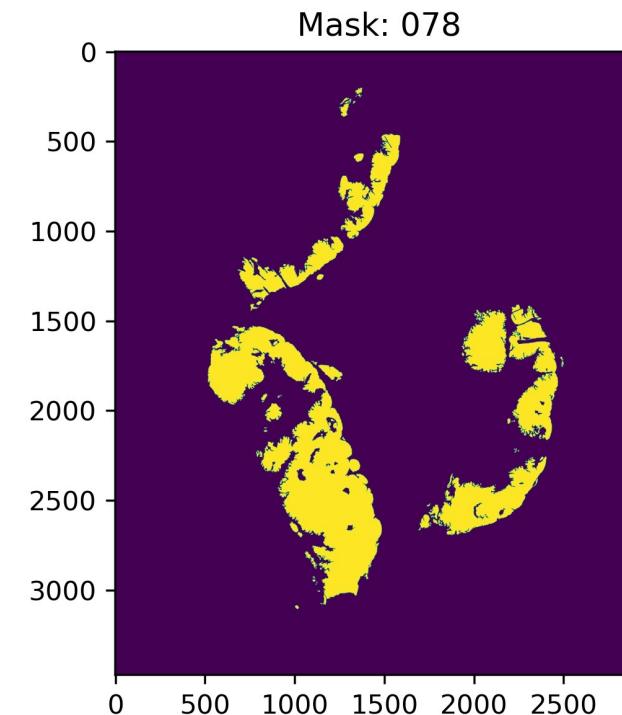
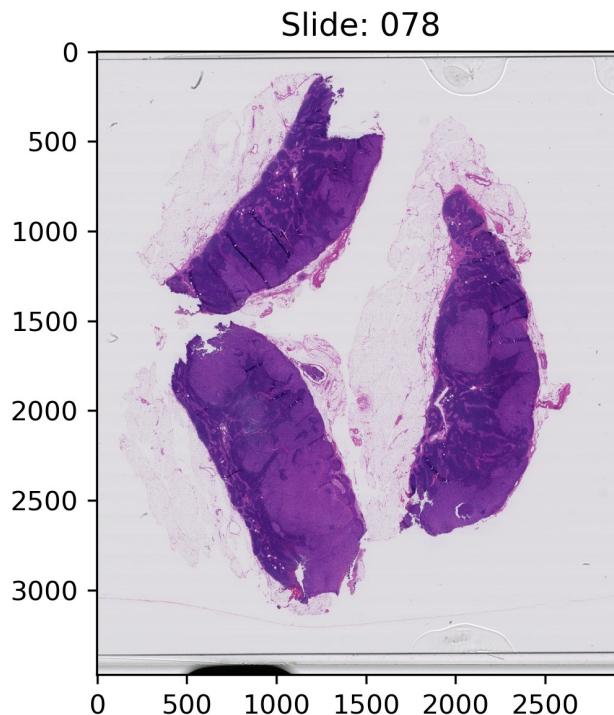
## Original Dataset: Gigapixel Pathology Images ([CAMELYON16 Dataset](#))

- Number of slides: 22
- Number of masks: 21 (there's no mask for slide 038)



# 1.2 Split the Dataset & Sanity Check

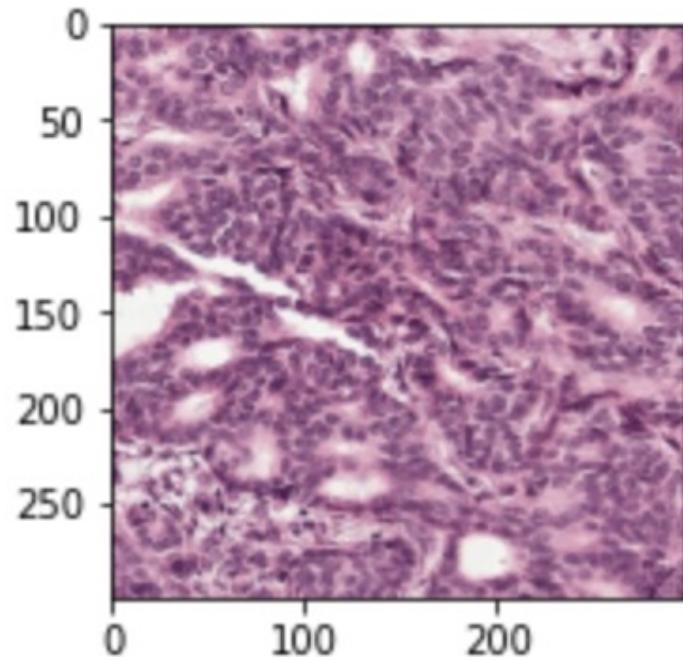
- Tried to use all available slides at first, the result was not quite good, and the training process took a long time.
- Manually select slides for training, validation, and test sets
- Check whether all slides contain cancer cells



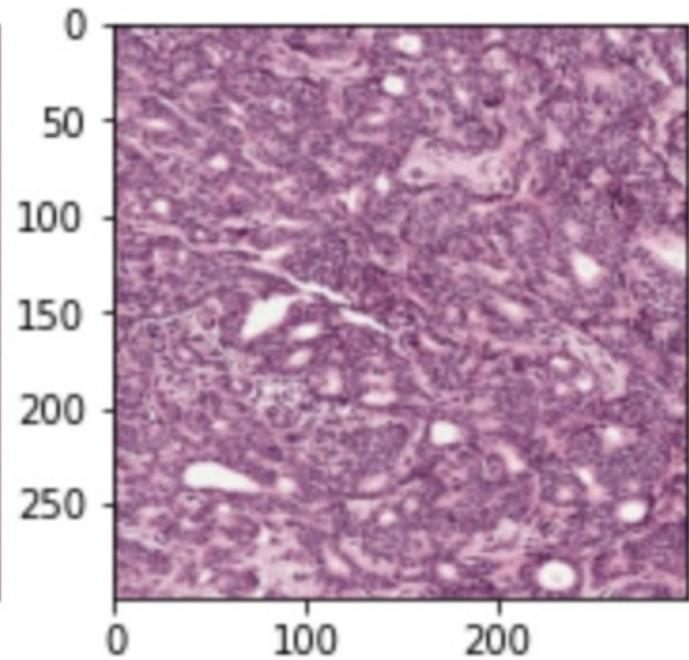
# 1.3 Generate Patches

## Strategies Used:

- Fixed patch size (299 x 299) at two different zoom levels
- Aligned at the same center point
- Tissue percentage  $\geq 30\%$
- Use the label of the lower zoom level (higher resolution)
- If there is any tumor pixel existing in the patch, then the label is 1 (positive)



Zoom Level 2



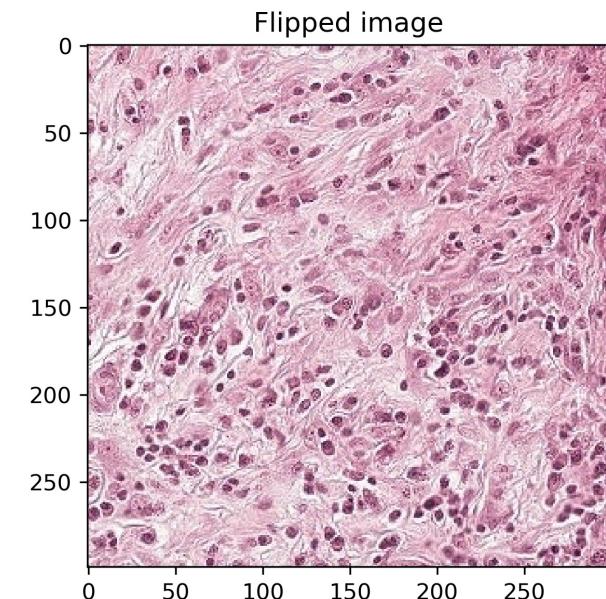
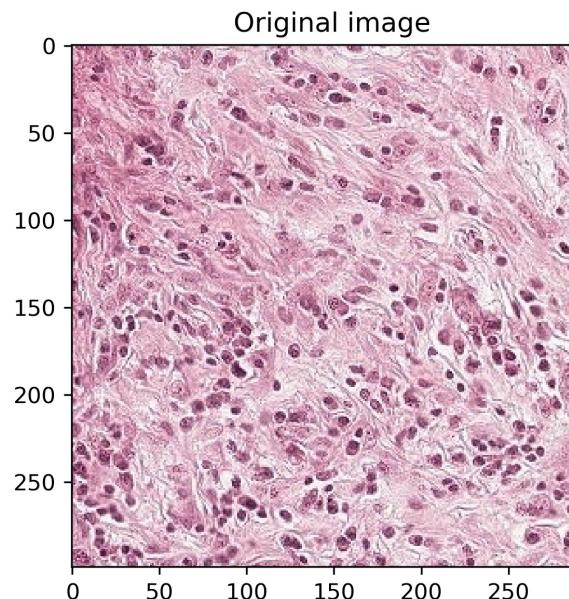
Zoom Level 3

# 1.4 Deal with Imbalanced Data

- ❖ Problem: More negative samples than positive samples.

## Strategies Tried:

- Data Augmentation
- Reduce the number of non-tumor samples, which has much more samples than tumor samples
- Adjust class weights at the training process
- Define custom loss functions to penalize undesired behaviors

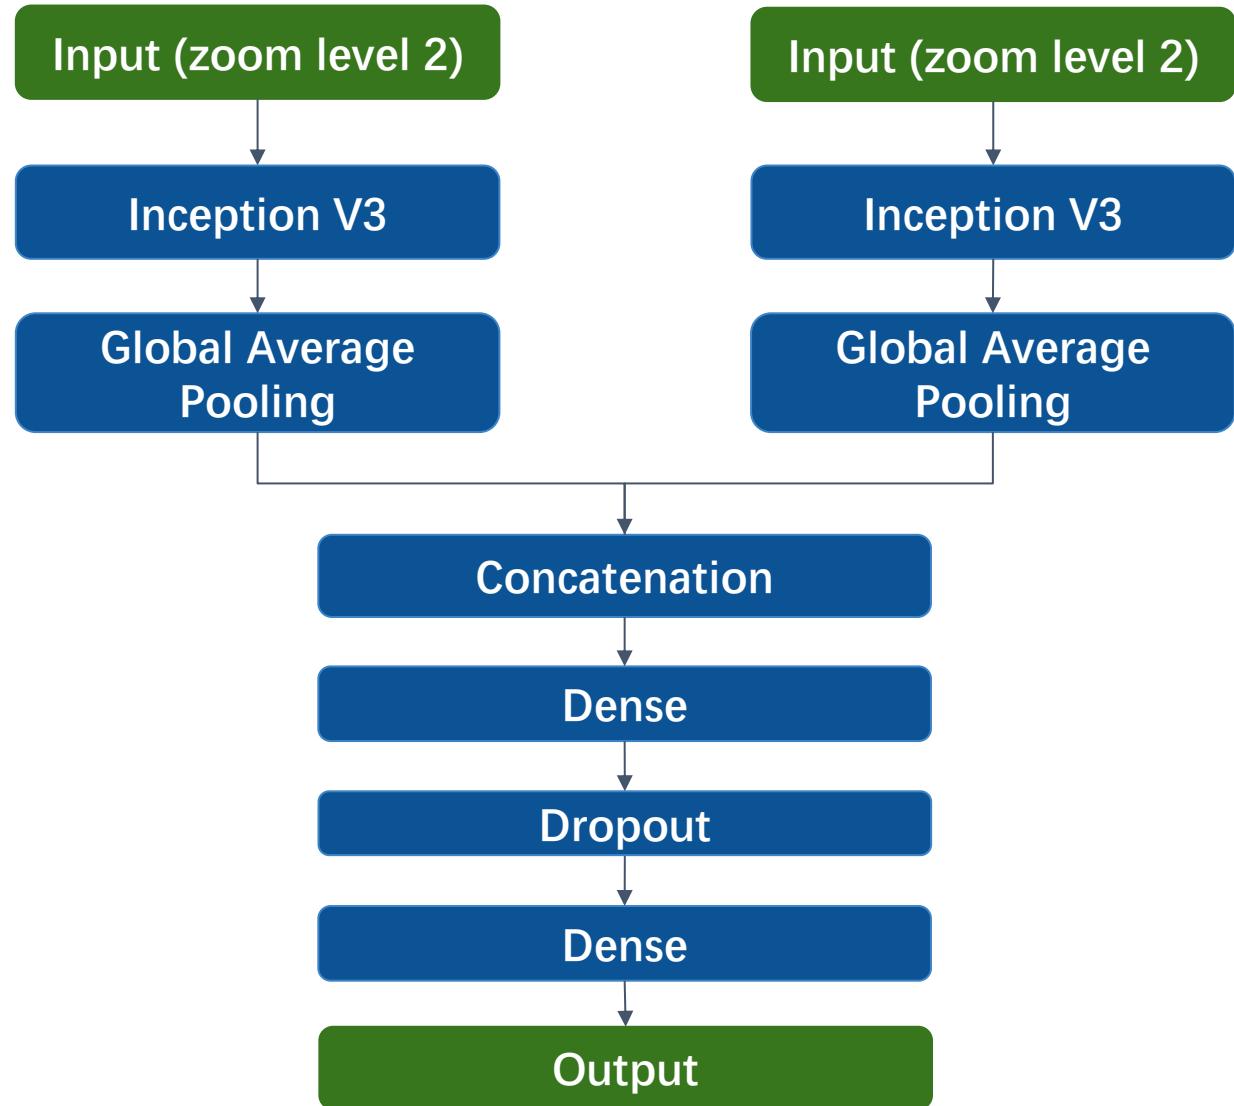


## II. Model Building

# 2.1 Model Building

## ➤ Features:

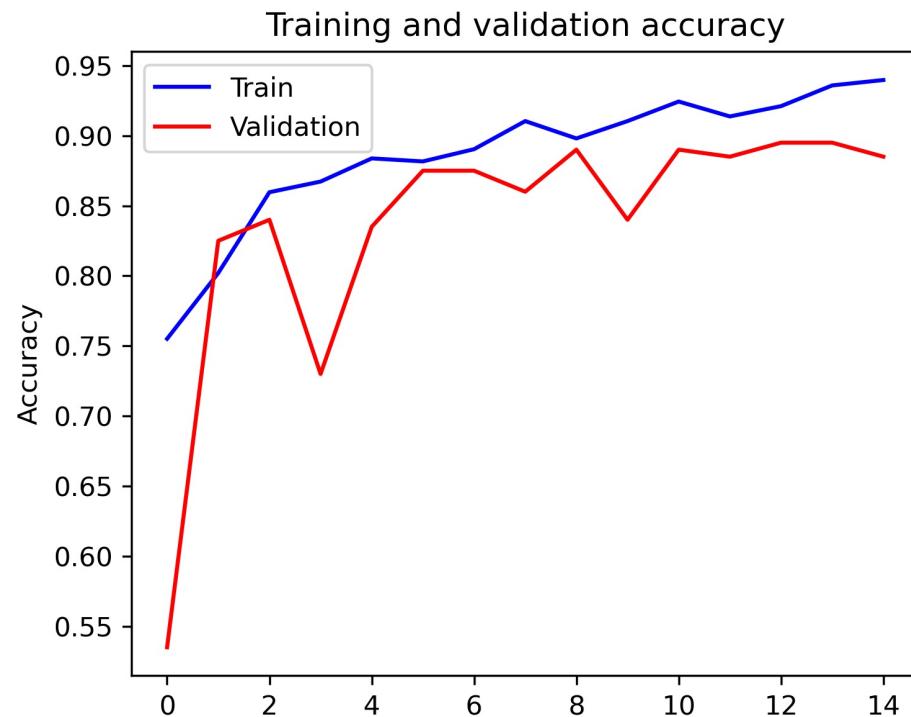
- Multi-input model
- Pretrained Inception V3 model on ImageNet (frozen)
- Global Average Pooling layers to reduce parameters
- Concatenation layer to leverage multi-scale inputs
- Dropout layer to mitigate overfitting



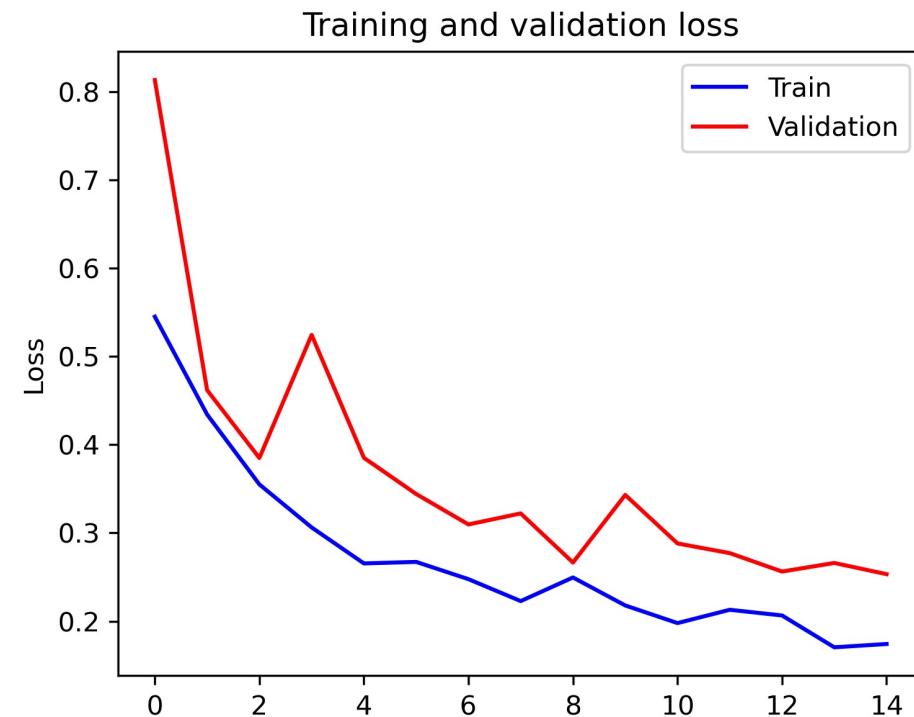
# III. Model Training and Evaluation

# 3.1 Loss and Accuracy

- Final training accuracy: 0.9448
- Final validation accuracy: 0.8850



- Final training loss: 0.1556
- Final validation loss: 0.2533



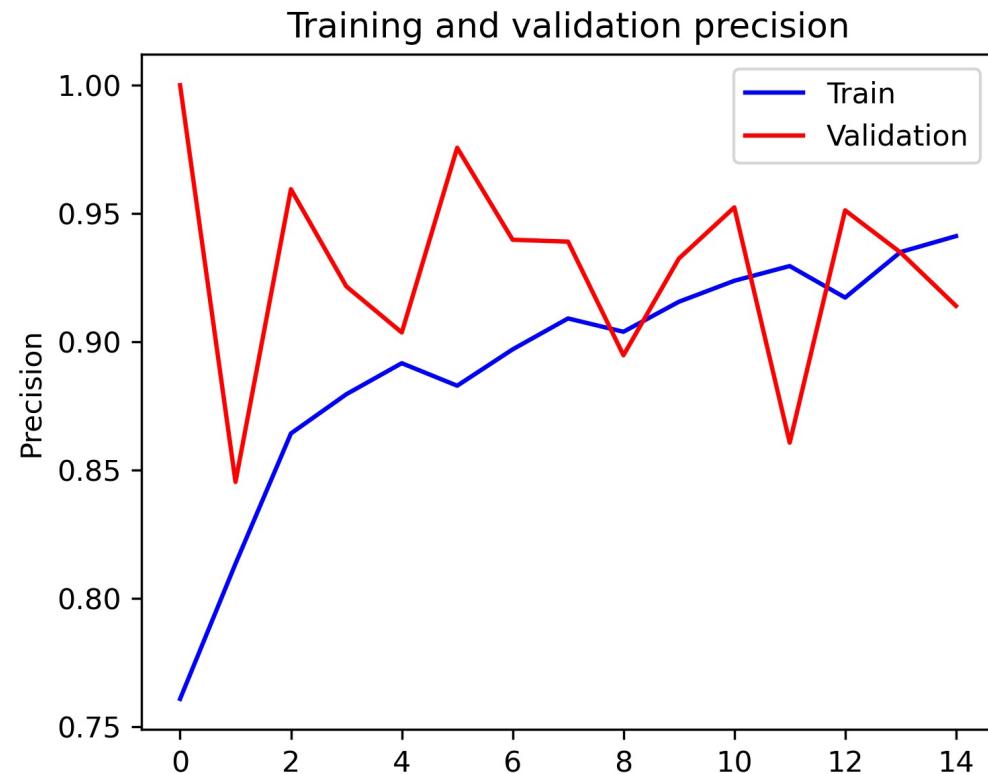
## 3.2 Precision

➤ **Precision:**

The percentage of predicted positives that were correctly classified

$$\text{precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

- Final training precision: 0.9551
- Final validation precision: 0.9140



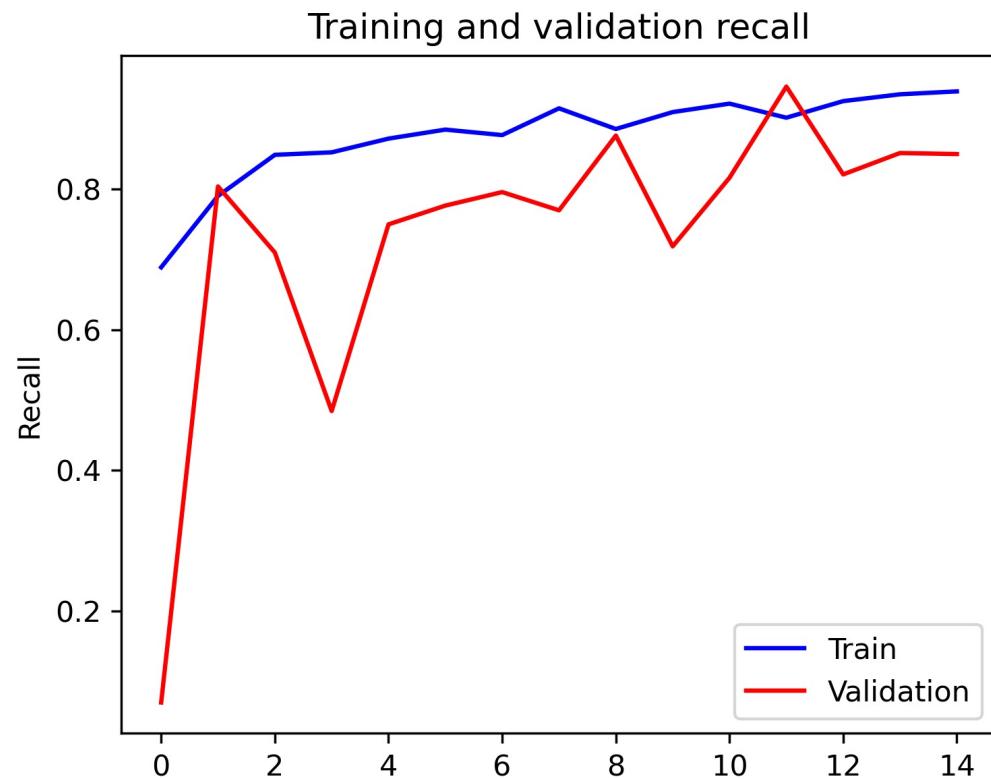
## 3.3 Recall

➤ **Recall:**

The percentage of actual positives that were correctly classified

$$\text{recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

- Final training recall: 0.9318
- Final validation recall: 0.8500



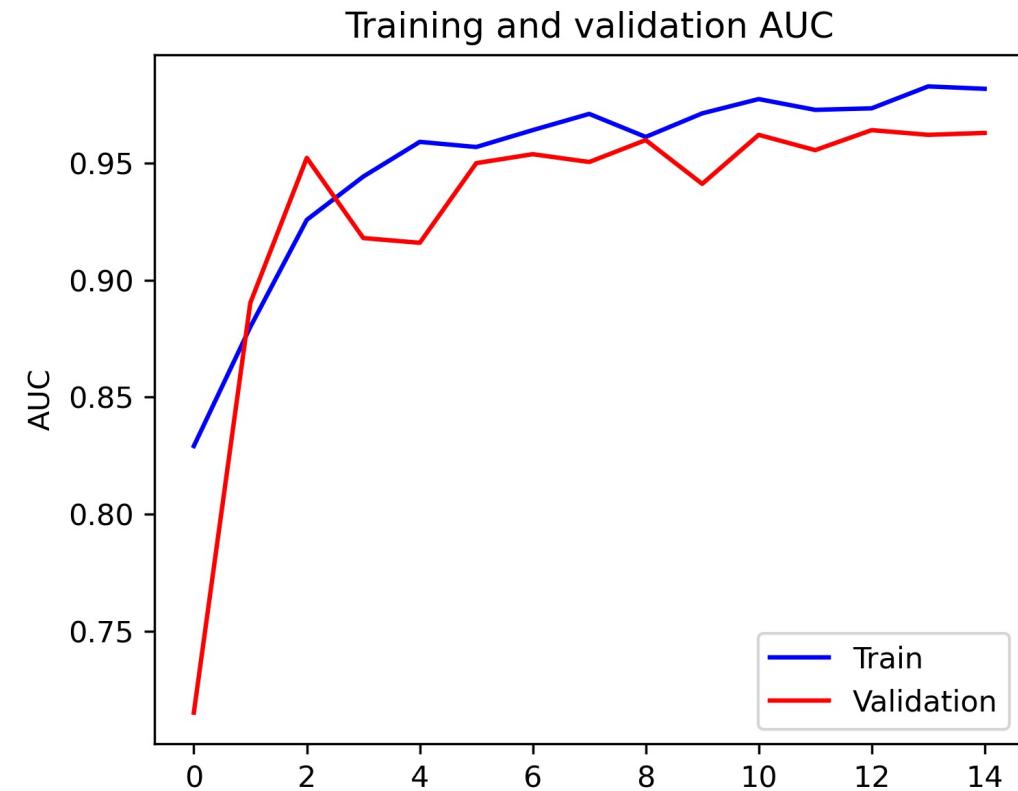
# 3.4 AUC (Area Under Curve)

## ➤ AUC:

AUC refers to the Area Under the Curve of a Receiver Operating Characteristic curve (ROC-AUC).

AUC ranges in value from 0 to 1.

- Final training AUC: 0.9862
- Final validation AUC: 0.9628

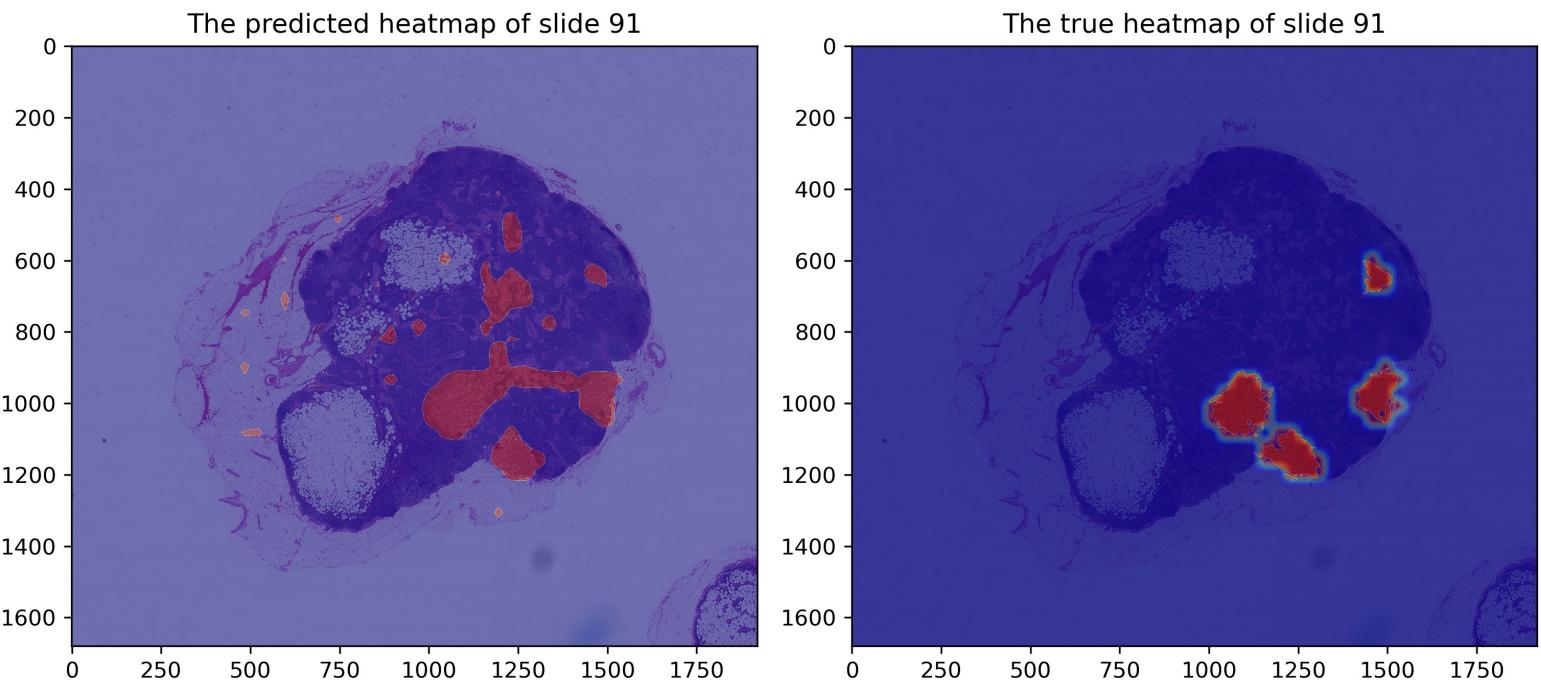


# IV. Model Predictions

# 4.1 Prediction Heatmaps

## Evaluation Metrics:

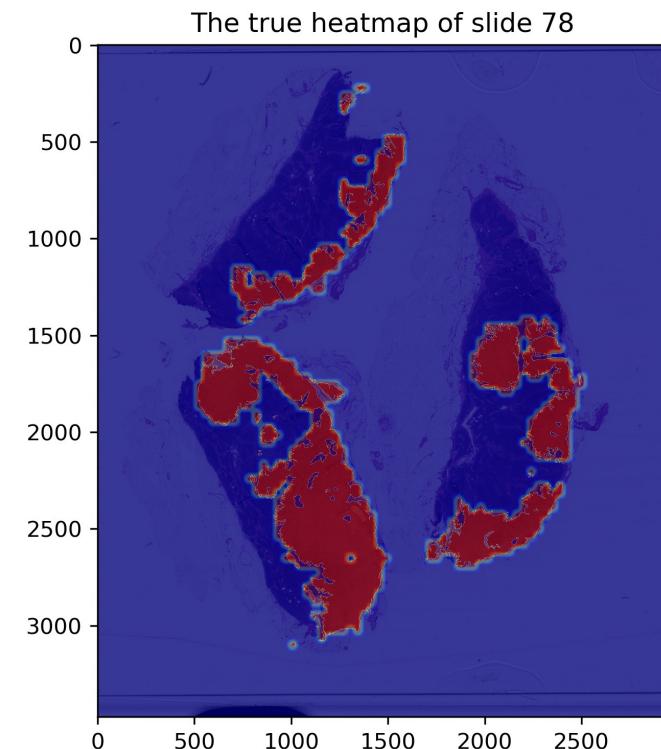
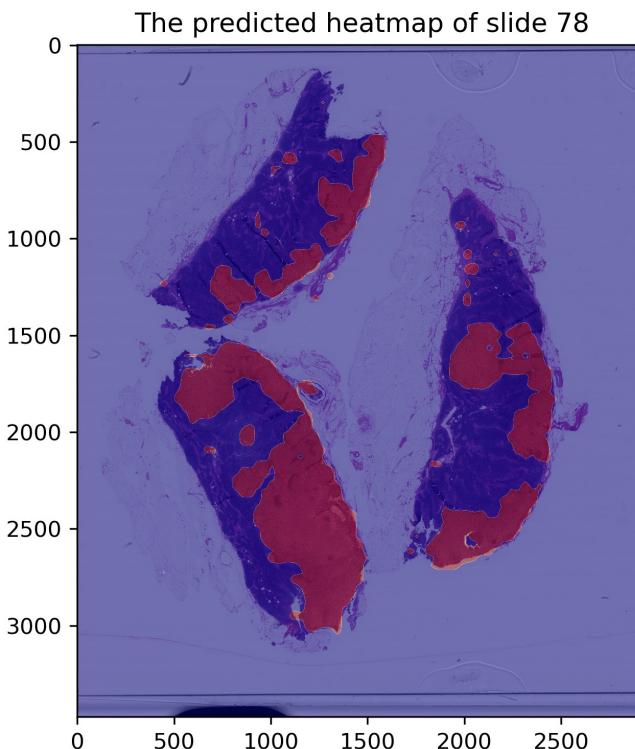
- AUC: 0.9508
- Precision: 0.4368
- Recall: 0.7451



# 4.1 Prediction Heatmaps

## Evaluation Metrics:

- AUC: 0.9811
- Precision: 0.9424
- Recall: 0.8064



# Conclusions

# Conclusions

In general, the model performs well, to improve it:

- Use lower level (higher resolution) slides
- Generate more patches for training (this may require more hardware resources)
- Configure the dataset for performance (cache, prefetch)
- Explore other model structures

# Thank you!

Weiwei Zhan, Haokai Zhao